Tanmay Bhuskute

+1 (732) 858-3027 | tanmay.v.bhuskute@gmail.com | github.com/Tanmayb05 | linkedin.com/in/tanmay-bhuskute

SUMMARY

Software Engineer and M.S. Computer Science candidate (Dec 2026) building large-scale ML pipelines, RAG systems, and data-quality/entity-resolution workflows with Python, PyTorch, PySpark, and cloud infrastructure. Current research in LLM-powered automation and evaluation.

EDUCATION

Master of Science in Computer Science Aug 2024 - Dec 2026
University of Cincinnati
Relevant Coursework: Advanced Algorithms, Deep Learning, Intelligent Data Analysis, Database Theory.

Bachelor of Engineering in Computer Engineering Aug 2018 - Jul 2022
AISSMS College of Engineering
Relevant Coursework: Data Mining, Data Analytics, Machine Learning, Object-Oriented Programming.

TECHNICAL SKILLS

ML & Data: PyTorch, TensorFlow, scikit-learn, PySpark, Delta Lake, LangChain, pandas, NumPy, Vertex AI.

Programming: Python, Java, Go, JavaScript, TypeScript, C/C++, SQL.

Web & Backend: React, Node.js, Django, Flask, FastAPI, REST APIs.

Databases: PostgreSQL, MySQL, Redis.

Developer Tools: Git, GitLab CI/CD, Jenkins, AWS, Docker, Kubernetes, ArgoCD, Helm, Terraform, Linux, Bash, SonarQube, Azure, Google Cloud Platform.

WORK EXPERIENCE

Graduate Researcher —- University of Cincinnati Feb 2026 – Present

- Designed two-stage memory-augmented LLM pipeline decoupling video understanding from device automation for Android GUI bug reproduction, evaluated against **ViBR** (GPT-4o + GroundingDINO + CLIP baseline) across 53 runs on 20+ apps.
- Demonstrated **25.7x latency reduction** (39.4s vs 1012.1s avg per run) while preserving task-completion fidelity, showing memory-reuse cuts redundant model inference at scale.

Software Development Engineer —- Siemens Industry Software Jul 2022 – Jun 2024

- Automated **Teamcenter X** infrastructure workflows with Python (Boto3), replacing legacy Ruby scripts and reducing deployment time by **40%** across AWS environments.
- Executed **44 production environment upgrades** in 4 months, diagnosing live infrastructure issues across Linux, EC2, databases, and Teamcenter X services.
- Designed **Ansible-based Blue-Green deployment** framework, reducing production upgrade downtime by **50%**.
- Migrated **40+ production environments** from process-based deployments to Kubernetes using ArgoCD, Helm, GitLab CI/CD, Amazon EFS, and HashiCorp Vault.
- Served as Quality Coach and SonarQube Specialist, increasing automated code coverage from **35% to 78%**; mentored **5 software engineers** through structured onboarding, driving engineering excellence.

Software Developer Intern —- Siemens Industry Software Mar 2022 – Jun 2022

- Automated provisioning of AWS resources with **Terraform**, reducing deployment time by **30%**; developed **Go + Terratest** framework catching **95%** of misconfigurations pre-production.

PROJECTS AND RESEARCH

NYC Mobility Forecasting [PySpark, Delta Lake, scikit-learn, FastAPI] Jul 2026

- Built end-to-end ML system forecasting NYC taxi/rideshare demand across **263 zones** (**2.3M-row** hourly dataset) via a Spark-based medallion pipeline; designed two-stage Poisson hurdle model improving MAPE **18.77% to 16.52%** and cutting RMSE **43%** over baseline Random Forest.

PolicyPilot – Medical Benefit Policy Intelligence Portal [React, TS, Node, Gemini, MCP] May 2026

- Architected evidence-grounded **RAG pipeline** over Google Gemini 2.5-flash with citation/version tracing to reconcile conflicting multi-source policy data and eliminate hallucination in high-stakes eligibility decisions.
- Designed **17-tool MCP server** exposing structured LLM tool-calling for entity extraction, cross-source record comparison, and confidence-scored data-quality triage, consumable by external AI agents.

EXTRACURRICULAR ACTIVITIES

- **Finalist, RevolutionUC 2025 Hackathon** (300 participants, 31 schools) | **SAFe Practitioner 5.1** (Mar 2023).