Dear Hiring Committee,

I saw the job post on **Amazon's careers site** and thought I'd be a good fit for it.

I am a Master's student in Computer Science at the University of Cincinnati, graduating December 2026, looking for a role as a **Software Engineer, AI/ML** on the **Annapurna Labs** Distributed Training team at **Amazon**. With hands on experience automating large-scale AWS infrastructure and building PyTorch-based performance evaluation pipelines, I am excited about the opportunity to contribute to **AWS Neuron** and help push training throughput on Trainium.

Here are reasons that make me a great fit for the role:

- **Distributed systems and performance profiling:** As a Graduate Researcher, I built a PyTorch-based evaluation pipeline that profiled a two-stage LLM system end to end to isolate a per-step inference bottleneck, achieving a 25.7x latency reduction (39.4s vs 1012.1s avg per run) across 53 runs on 20+ Android apps while preserving task-completion fidelity.

- **Large-scale AWS infrastructure experience:** At Siemens Industry Software, I automated Teamcenter X infrastructure workflows with Python (Boto3), cutting deployment time by 40% across AWS environments, and executed 44 production environment upgrades in 4 months, diagnosing live infrastructure issues across Linux, EC2, and databases.

- **System design across the stack:** I migrated 40+ production environments from process-based deployments to Kubernetes using ArgoCD, Helm, and GitLab CI/CD, and designed an Ansible-based Blue-Green deployment framework that reduced production upgrade downtime by 50%, giving me practical experience coordinating changes across interdependent infrastructure layers.

- **Engineering quality and collaboration:** As Quality Coach and SonarQube Specialist for the Agile Release Train, I increased automated code coverage from 35% to 78% and mentored 5 software engineers, reflecting the cross-team collaboration this role requires when tracing performance gaps across frameworks, compiler, and runtime teams.

I am available for full-time work after my graduation in December 2026. As a Master's student at the University of Cincinnati and someone who values professionalism and growth, I am enthusiastic about contributing to Neuron's mission of accelerating deep learning and GenAI workloads on Trainium.

Please find my CV attached below. I look forward to hearing from you. I would be excited to work at the intersection of ML and distributed systems, helping shape the direction of AI acceleration technology.
Regards,
Tanmay Bhuskute,
bhuskutv@mail.uc.edu | tanmay.v.bhuskute@gmail.com | +1 (732) 858-3027,
3424 Brookline Ave, Flat 19, Cincinnati, OH 45220.