Dear Hiring Committee,

I saw the job post on **NVIDIA's careers website** and thought I'd be a good fit for it.

I am a Master's student in Computer Science at the University of Cincinnati, graduating December 2026, looking for a role as a **Software Engineer, Deep Learning Libraries** at **NVIDIA**. With hands on experience in C/C++, build automation, and scalable infrastructure, I am excited about the opportunity to contribute to **NVIDIA** and support the team building the deep learning libraries used across the world.

Here are reasons that make me a great fit for the role:

- **Strong C/C++ and object-oriented development background:** My coursework and projects are built on C/C++ and object-oriented design, and I have applied this foundation to production systems at Siemens Industry Software, where reliability and performance mattered at scale.

- **Experience building scalable automation for build, test, and release processes:** At Siemens, I automated Teamcenter X infrastructure workflows with Python (Boto3), replacing legacy scripts and cutting deployment time by 40% across AWS environments, and executed 44 production environment upgrades in 4 months.

- **Hands on experience with industry-standard tools across platforms:** I migrated 40+ production environments to Kubernetes using ArgoCD, Helm, GitLab CI/CD, and HashiCorp Vault, and used Git, Jira, and Jenkins-style CI/CD pipelines daily to maintain and test environments across new configurations.

- **Pragmatic, "it just works" approach to automation and collaboration:** As Quality Coach and SonarQube Specialist, I increased automated code coverage from 35% to 78% by enforcing coding standards and quality gates, and mentored 5 software engineers through structured onboarding, reflecting my passion for enabling teammates through automation.

- **Ability to work independently and apply deep learning libraries:** My graduate research builds a memory-augmented LLM pipeline using PyTorch-based models, Gemini, and Vertex AI for Android bug reproduction, achieving a 25.7x latency improvement, strengthening my ability to define project scope and deliver independently.

I am available for full-time work after December 2026. As a Master's student at the University of Cincinnati and someone who values professionalism and growth, I am enthusiastic about contributing to NVIDIA's Deep Learning Libraries team.

Please find my CV attached below. I look forward to hearing from you. I would be excited to help build the software that powers breakthroughs across image classification, speech recognition, and natural language processing worldwide.

Regards,

Tanmay Bhuskute,
bhuskutv@mail.uc.edu | tanmay.v.bhuskute@gmail.com | +1 7328583027,
3424 Brookline Ave, Flat 19, Cincinnati, OH 45220.