Dear Hiring Committee,

I saw the job post on the Sola careers page and thought I'd be a good fit for it.

I am a Master's student in Computer Science at the University of Cincinnati, graduating December 2026, looking for a **Software Engineering** role at **Sola**. With production experience spanning distributed backend infrastructure and hands-on AI research, I am excited about the opportunity to contribute to **Sola** and help build reliable, observable execution on top of non-deterministic AI models.

Here are reasons that make me a great fit for the role:

- **Owning AI features end-to-end, from experimentation to production:** My graduate research designs and evaluates a two-stage memory-augmented LLM pipeline for Android automation, taking it from an early prototype through a structured evaluation against a published baseline (GPT-4o + GroundingDINO + CLIP), achieving a 25.7x latency reduction while preserving task fidelity across 20+ real-world apps.

- **Building and operating reliable distributed backend systems:** At Siemens Industry Software, I migrated 40+ production environments to Kubernetes using ArgoCD, Helm, and GitLab CI/CD, and designed an Ansible based Blue Green deployment framework that cut upgrade downtime 50% with automated rollback within 15 minutes.

- **Proficiency in Python and modern ML frameworks:** I am proficient in Python and TypeScript, with hands-on experience across PyTorch, scikit-learn, and Vertex AI, including building and benchmarking a PyTorch LSTM classifier on 400K+ headlines and provider-agnostic LLM evaluation harnesses spanning Gemini, GPT-4o, Llama, and Qwen.

- **Building observability into production systems:** At Siemens, I diagnosed live production issues across 40+ environments using CloudWatch and Datadog log correlation, and built a self-service developer platform adopted by 15 engineers, the same instinct for monitoring and measuring system health I would bring to evals and observability infrastructure for AI systems in production.

- **Comfort with ambiguity at an early stage:** I have shipped systems end to end in production, from Terraform-based AWS provisioning to a self-service CLI and dashboard, and enjoy the ambiguity of figuring out what is feasible to build rather than working from a fixed spec. A small, in-person team turning cutting-edge AI research into shipped features is exactly the environment I want to be in.

I am available to work with you after my graduation in December 2026, and am glad to relocate to New York City. As a Master's student at the University of Cincinnati and someone who enjoys turning AI research into reliable production systems, I am enthusiastic about contributing to Sola and helping bring state-of-the-art agentic capabilities to production at scale.

Please find my CV attached below. I look forward to hearing from you. Sola's mission of building reliable, deterministic execution on top of non-deterministic AI models is exactly the kind of hard, high-impact problem I want to help solve.

Regards,

Tanmay Bhuskute,
bhuskutv@mail.uc.edu | tanmay.v.bhuskute@gmail.com | +1 (732) 858-3027,
3424 Brookline Ave, Flat 19, Cincinnati, OH 45220.