Dear Hiring Committee,

I saw the job post on **email outreach from your recruiting team** and thought I'd be a good fit for it.

I am a Master's student in Computer Science at the University of Cincinnati, graduating December 2026, looking for a role as a **Senior Gen AI/Python Engineer** at **Synchrony Systems**. My work sits exactly where this role does: production Python, LLM powered agents, structured extraction, and FastAPI services that have to be correct rather than merely impressive in a demo.

Here are reasons that make me a great fit for the role:

- **Production grade Python with real engineering discipline:** At Siemens Industry Software I wrote Python and Boto3 automation that replaced legacy Ruby services and cut deployment time by 40% across AWS environments, shipped through code review, Docker, and GitLab CI/CD. As Quality Coach and SonarQube Specialist I raised automated test coverage from 35% to 78% by enforcing coding standards and CI quality gates.

- **Building LLM agents and multi step orchestration:** My graduate research is a two stage agentic pipeline on Gemini 2.5 Pro and Vertex AI that extracts a structured task memory once, then conditions a closed loop automation agent on it instead of re running inference at every step. It delivered a 25.7x latency reduction, 39.4 seconds versus 1012.1 seconds per run, against the published GPT-4o based ViBR baseline while preserving task completion fidelity across 53 runs on more than 20 applications.

- **Structured output, RAG, and confidence based disambiguation:** In PolicyPilot I architected an evidence grounded RAG pipeline over Gemini 2.5-flash that forces citation backed answers, and designed a 17 tool MCP server for LLM tool calling with schema validation at every data boundary. I added a confidence scoring layer using evidence count, negation detection, and subjectivity flags to route uncertain cases to human review, chosen after benchmarking hallucination rates across models.

- **FastAPI, async orchestration, and large analytical datasets:** I built a FastAPI prediction microservice backed by PostgreSQL and Redis caching, containerized with Docker, serving a model trained on a 2.3 million row Spark and Delta Lake dataset across 263 zones. I also built a provider agnostic model abstraction and a config driven evaluation harness so prompt and model changes could be validated through reproducible regression runs.

I am available for full time work after December 2026, matching my graduation date. As a Master's student at the University of Cincinnati and someone who values professionalism and growth, I am enthusiastic about contributing to **Synchrony Systems** and helping ship reliable, well evaluated Gen AI capabilities into production.

Please find my CV attached below. I look forward to hearing from you. **Building agentic systems that stay accurate under real production data is the problem I most want to work on, and this role is squarely that work**.

Regards,

Tanmay Bhuskute,
bhuskutv@mail.uc.edu | tanmay.v.bhuskute@gmail.com | +1 (732) 858-3027,
3424 Brookline Ave, Flat 19, Cincinnati, OH 45220.